

The COIN-OR Open Solver Interface 2.0 Redux (Zombies Attack!)

Lou Hafer¹ Matthew Saltzman²

¹Department of Computer Science
Simon Fraser University

²Department of Mathematical Sciences
Clemson University.

INFORMS Annual Meeting
Austin, Texas
November 10, 2010

What is OSI?

- ▶ A cross-solver API
- ▶ Lower level than most solver APIs
 - ▶ Instance management
 - ▶ Algorithm control (e.g., pivot-level simplex) is a goal (honored more often than not in the breach)
 - ▶ Intended as a “crossbar switch” to connect applications to solvers
- ▶ One of the original COIN-OR projects (a product of impetuous youth and inexperience)

A Fractured History

- ▶ Original goal for OSI: “solver independent”
 - ▶ Too ambitious—can’t reproduce outcomes reliably with different solvers

A Fractured History

- ▶ Original goal for OSI: “solver independent”
 - ▶ Too ambitious—can’t reproduce outcomes reliably with different solvers
- ▶ More achievable: “solver agnostic”
 - ▶ Common way to interact with solvers, but don’t enforce solver behavior
 - ▶ Success (more or less) at source code level

Source Level Control

- ▶ `OsiXxxSolverInterface` derived from `OsiSolverInterface` for various values of `Xxx` ($Xxx \in \{\text{Clp}, \text{Cpx}, \text{Grb}, \dots\}$).
- ▶ User program instantiates concrete `OsiXxxSolverInterface` object
- ▶ So `Xxx` needs to be specified in source
- ▶ Selection can be controlled by compiler directives (`#ifdef`)
- ▶ Specified solver engine library required at link time

Source Level Control

- ▶ `OsiXxxSolverInterface` derived from `OsiSolverInterface` for various values of `Xxx` ($Xxx \in \{\text{Clp}, \text{Cpx}, \text{Grb}, \dots\}$).
- ▶ User program instantiates concrete `OsiXxxSolverInterface` object
- ▶ So `Xxx` needs to be specified in source
- ▶ Selection can be controlled by compiler directives (`#ifdef`)
- ▶ Specified solver engine library required at link time
- ▶ Solver engine library must ship with user's binary!
 - ▶ For example, our CBC binary must ship with CLP as LP solver, because we can't ship other LP solvers.
 - ▶ An end user who wants another solver engine must build from source

Shared Library for Solver Engine

- ▶ If solver libraries are shared (`libxxx.so` on Unix, `DLL` on Windows)
 - ▶ We can ship without solver libs—they are linked at load time
 - ▶ Still need to specify solver in source code
 - ▶ User still needs to have solver libs
 - ▶ CPLEX isn't shipped as shared lib!

Dynamic Loading of Solver Engine

- ▶ `dlopen()`, `dlsym()`, etc., in Linux, other calls in Windows and other Unix systems
- ▶ There are cross-platform libraries for this task (GNOME glib, GNU libtool)
- ▶ Solver Engine loaded at runtime, not needed at link time
- ▶ But still tied to solver in source

Plugins: An Alternate Definition of “Solver Independence”

- ▶ Delay decision of what solver to instantiate until runtime
 - ▶ User declares abstract interface object (`OsiSolverInterface`)
 - ▶ Asks factory object to create a specified concrete implementation
- ▶ Solver engine loaded dynamically when implementation object is created
- ▶ Now we can ship CBC (say) and let the end user decide at runtime which LP engine to use
- ▶ Plugin builder and user need access to solver engine

Plugins: An Alternate Definition of “Solver Independence”

- ▶ Delay decision of what solver to instantiate until runtime
 - ▶ User declares abstract interface object (`OsiSolverInterface`)
 - ▶ Asks factory object to create a specified concrete implementation
- ▶ Solver engine loaded dynamically when implementation object is created
- ▶ Now we can ship CBC (say) and let the end user decide at runtime which LP engine to use
- ▶ Plugin builder and user need access to solver engine
- ▶ Now we're breaking source code compatibility. . .

Plugins: An Alternate Definition of “Solver Independence”

- ▶ Delay decision of what solver to instantiate until runtime
 - ▶ User declares abstract interface object (`OsiSolverInterface`)
 - ▶ Asks factory object to create a specified concrete implementation
- ▶ Solver engine loaded dynamically when implementation object is created
- ▶ Now we can ship CBC (say) and let the end user decide at runtime which LP engine to use
- ▶ Plugin builder and user need access to solver engine
- ▶ Now we're breaking source code compatibility. . .
- ▶ . . . so we can think about what we would do if we were starting with a clean slate

What Else is Wrong?

- ▶ Front and back ends can get out of sync
- ▶ Interface changes break everything
- ▶ Extensions are difficult
- ▶ Feature-complete new shims are painful to implement
- ▶ No upgrade path
- ▶ Too many tasks are implemented in the shim layer (e.g., caching)—no way to implement common code
- ▶ No way for caller to know what capabilities are available or missing
- ▶ ...

What do we want?

- ▶ Writing shims should be straightforward (not much harder than other APIs)
- ▶ Using the interface should be straightforward (not much harder than using an unwrapped solver)
- ▶ Performance penalty should be minimal
- ▶ The interface should provide a useful set of capabilities
- ▶ The interface should be extensible
 - ▶ New capabilities should be easy to offer through the interface
 - ▶ Hooking the solver directly should truly be a last resort

Decouple Interface from Implementation

- ▶ In C++, this is a matter of programmer discipline
- ▶ Necessary to implement plugins
- ▶ Necessary to implement extensions

Decouple Interface from Implementation

- ▶ In C++, this is a matter of programmer discipline
- ▶ Necessary to implement plugins
- ▶ Necessary to implement extensions
- ▶ Abstract base class exposes only public interface—defines module semantics
- ▶ Concrete implementation object derived from abstract base class
- ▶ User asks factory method to return concrete object
- ▶ All implementation details and private data are hidden

Modularization

- ▶ Clusters of related methods to accomplish tasks
 - ▶ Core instance management
 - ▶ Presolve
 - ▶ Linear algebra, basis management
 - ▶ Simplex
 - ▶ B&C control
 - ▶ ...
- ▶ Capabilities managed via a map (name, version, factory)
 - ▶ Name defines semantics via abstract base class
 - ▶ Loading an interface returns a pointer to a concrete object
 - ▶ Upgrades (for developer)
 - ▶ Fallbacks (for user)
- ▶ Modules associated with an instance need to match solver engine

Inter-Module Communication

- ▶ Modules are users too
 - ▶ Need access to capabilities
- ▶ Incoming module responsible for replacing existing module with the same functionality
 - ▶ Extract data
 - ▶ Replace capabilities
 - ▶ Unload old module

Callbacks

- ▶ C engine callback handled by registering a function with specific signature
- ▶ C++ engine callback is method derived from virtual base method
- ▶ Different engines define different categories of callbacks, provide different types of information, and allow different sets of actions
- ▶ OSI could implement a limited set of callback actions (check for abort flag and abort) in a common set of hooks
- ▶ Much more than that requires exposing solver-specific interfaces

Parameter Management

- ▶ Infinite variety
- ▶ Might be able to identify a set of common ones
- ▶ Map/table?
- ▶ Probably need to expose solver-specific interface for less common settings

Other matters

- ▶ Message handling
- ▶ Interactions with other COIN-OR components
- ▶ ???

Lawyers, Guns, and Money

- ▶ Plugins mitigate license compatibility issues
- ▶ GPL requires any program that “includes” GPL code must be GPL
- ▶ But plugins are not “included” in programs that use them

Questions? Suggestions?